



Policy for the responsible use of AI in government

24-0180

Version 2.0



Digital Transformation Agency



© Commonwealth of Australia (Digital Transformation Agency) 2025 | Version 1801

With the exception of the Commonwealth Coat of Arms and where otherwise noted, this product is provided under a Creative Commons Attribution 4.0 International Licence. (<http://creativecommons.org/licenses/by/4.0/legalcode>)

The Digital Transformation Agency has tried to make the information in this product as accurate as possible. However, it does not guarantee that the information is totally accurate or complete. Therefore, you should not solely rely on this information when making a commercial decision.

Digital Transformation Agency is committed to providing web accessible content wherever possible.

If you are having difficulties with accessing this document, please email: ai@dta.gov.au.

Version: 2.0

Contents

Policy introduction	4
Policy aims	5
Embrace the opportunity	5
Strengthen public trust	5
Adapt to change	5
Implementation	6
Application	6
National security carveouts	6
Existing frameworks	7
Definitions	7
Timeframes	7
In-scope AI use cases	8
Principles and requirements	9
Strategy and oversight	10
Principles	10
Mandatory requirements	10
Preparedness and operations	12
Principles	12
Mandatory requirements	12
Recommended Actions	13
Agencies should consider	13
AI use case impact assessment	14
Principles	14
Mandatory requirements	14
Appendix A: Related frameworks for AI	17
Artificial intelligence	17
Automated decision-making	17
Data and digital	17
Appendix B: Definitions	18
Artificial intelligence	18
AI use case	19
AI incident	20
Appendix C: In-scope AI use cases	21
Criteria and areas of consideration	21
Experimentation	22

Policy introduction

The increasing adoption of artificial intelligence (AI) is reshaping the economy, society and government. Government agencies are embracing the AI opportunity to boost efficiency, enable better data analysis and evidence-based decisions, and improve service delivery for people and business.

The Australian Government is committed to capturing the opportunities of AI in a way that facilitates innovation while mitigating harms. Government has an elevated responsibility to adopt AI safely and should be held to a high ethical standard of use.

The December 2025 update to the policy strengthens government's approach to safe and responsible AI through new measures on AI governance. It introduces requirements for agencies to develop a strategic approach to adopting AI, establish an approach to operationalise the responsible use of AI, ensure designated accountability for AI use cases, and undertake risk-based use case-level actions.

In recognition of the speed and scale of change in this area, the policy will evolve over time as the technology changes, leading practices emerge, the broader regulatory environment develops, and government's AI maturity improves.

Policy overview

This policy provides a framework to enable the accelerated and sustainable adoption of AI by agencies. It allows them to embrace AI while continuing to position government as an exemplar in responsible AI use, in line with the Australian community's expectations.

The policy aims to create a coordinated approach to government's use of AI and has been designed to complement and strengthen – not duplicate – existing frameworks in use by the APS.

The policy outlines principles and requirements under the following sections: **strategy and oversight, preparedness and operations**, and **AI use case impact assessment**.

Each section of the policy is underpinned by principles and outlines the mandatory requirements agencies must follow. Recommended actions and considerations are also included where relevant.

Policy aims

This policy aims to ensure that government plays a leadership role in embracing AI for the benefit of Australians while ensuring its safe, ethical and responsible use, in line with community expectations.

Embrace the opportunity

This policy aims to provide a unified approach to enable government to accelerate AI adoption and embrace the AI opportunity. It is designed to reduce barriers to government adoption by helping agencies confidently approach AI governance and implementation.

It aims to ensure agencies have the right settings in place to take advantage of the opportunities presented by AI and fully realise benefits such as increased and improved efficiency, accuracy and service delivery.

Strengthen public trust

This policy aims to strengthen public trust in government adoption of AI by positioning the Australian Government as an exemplar in safe and responsible AI use.

It is designed to enable the responsible use of AI across government, through setting consistent requirements for transparency and accountability, and by requiring risk-based oversight of AI use cases.

Adapt to change

This policy aims to embed a forward-leaning, adaptive approach for government's use of AI that evolves as the technological and policy environment changes.

It supports agencies at different stages of their AI adoption journey and sets requirements that scale with the agency's use of AI.

Implementation

Application

This version of the policy (v2.0) is effective **15 December 2025**. It replaces version v1.1 of the policy which came into effect 1 September 2024.

All [non-corporate Commonwealth entities](#) (NCEs), as defined by the *Public Governance, Performance and Accountability Act 2013*, must apply this policy.

[Corporate Commonwealth entities](#) are also encouraged to apply this policy.

National security carveouts

This policy does not apply to:

- the use of AI in the defence portfolio.
- the 'national intelligence community' (NIC) as defined by Section 4 of the [Office of National Intelligence Act 2018](#).

The NIC includes:

- Office of National Intelligence (ONI)
- Australian Signals Directorate (ASD)
- Australian Security Intelligence Organisation (ASIO)
- Australian Secret Intelligence Service (ASIS)
- Australian Geospatial-Intelligence Organisation (AGO)
- Defence Intelligence Organisation (DIO)
- Australian Criminal Intelligence Commission (ACIC)
- the intelligence role and functions of the Australian Transaction Reports and Analysis Centre (AUSTRAC), Australian Federal Police (AFP), the Department of Home Affairs and the Department of Defence.

Defence and members of the NIC may voluntarily adopt elements of this policy where they are able to do so without compromising national security capabilities or interests.

Existing frameworks

The challenges raised by government use of AI are complex and inherently linked with other considerations, such as the APS Code of Conduct, data governance, cyber security, privacy and ethics practices.

This policy has been designed to complement and strengthen – not duplicate – existing frameworks, legislation and practices that touch on government's use of AI.

This policy must be read and applied alongside existing frameworks and laws to ensure agencies meet all their obligations.

Definitions

For the purposes of this policy, agencies should apply the definition of AI provided by the Organisation for Economic Co-operation and Development (OECD) and the following definition of AI use case:

An AI use case is a specific application of an AI system or systems to achieve certain objectives or perform certain tasks.

Definitions and how to apply them – including an optional approach to group AI use cases for some general-purpose AI solutions – is available in **Appendix B**. The appendix also defines an AI incident.

Timeframes

This policy provides implementation timeframes for agencies to meet some of its requirements. While agencies may need this time to action requirements, agencies should implement them sooner if practicable. Agencies could consider putting in place interim processes and building out their approach as they reach the specified implementation deadline.

In-scope AI use cases

This policy specifies actions that apply at the use case level. AI use cases in scope of this policy (referred to as in-scope AI use cases) are use cases that meet any criteria in **Appendix C**.

In addition to the criteria, the appendix lists areas of AI use to consider that are not automatically high risk, but are more likely to involve risks that require careful attention through an impact assessment. It also provides information on how to apply the policy for agencies experimenting with AI.

Principles and requirements

The following sections outline the policy principles and requirements: strategy and oversight, preparedness and operations, and AI use case impact assessment.




Strategy and oversight



Preparedness and operations



AI use case impact assessment





Strategy and oversight

The principles and requirements included in this section are designed to enable a forward-leaning approach to agency AI adoption. These requirements set out AI accountability at the agency level and use case level and seek to build trust through transparency.

Principles

- Adopt AI to enhance efficiency, decision-making, policy outcomes and government service delivery for the benefit of Australians.
- Have clear accountabilities for the adoption of AI and understand its use.
- Build public trust through transparency about government AI use.

Mandatory requirements

AI transparency statement

Agencies **must** make a publicly available statement outlining their approach to AI adoption and use, as prescribed under the [Standard for transparency statements](#).

The statement **must** be reviewed and updated annually or sooner, should the agency make significant changes to its approach to AI.

Agencies **must** notify the DTA when they publish and make any changes to their AI transparency statement by emailing ai@dtg.gov.au.

Strategic position on AI adoption

Agencies **must** develop a strategic position on AI adoption within 6 months of this policy taking effect. This position is to emphasise how AI opportunities can be identified and embraced by the agency.

Agencies **must** communicate their strategic position on AI to give staff clear direction on AI adoption. In line with their current and anticipated use of AI, agencies can develop a standalone AI strategy, augment an existing strategy or create other materials to communicate the approach to staff.

Accountable officials

Agencies **must** designate accountable official(s) to take accountability for implementing this policy.

Agencies **must** follow the [Standard for accountability](#) when designating accountable official(s) and implementing this requirement. The responsibilities of accountable officials are set in the standard.

Agencies **must** notify the DTA when they designate and make any changes to their accountable official(s) by emailing ai@dtg.gov.au.

Accountable use case owners

Agencies **must** designate an accountable use case owner for each in-scope AI use case within 12 months of this policy taking effect. Accountable official(s) are to maintain a register of accountable use case owners.

Agencies **must** follow the [Standard for accountability](#) when implementing this requirement. The responsibilities of accountable use case owners are set in the standard.

Internal AI use case register

Agencies **must** create a register of in-scope AI use cases to enable accountable official(s) to record accountable use case owners within 12 months of this policy taking effect.

Agencies **must** share the register with the DTA every 6 months, commencing from when they create the register to meet the above requirement.

The [Standard for accountability](#) lists the minimum fields agencies must capture in the use case register. Agencies can add additional fields to meet their organisational needs. An existing register may be reused for the purposes of meeting this requirement. The standard also provides the instructions for how to share agency registers with the DTA.



Preparedness and operations

The principles and requirements included in this section standardise key elements of AI governance that allow agencies to build AI capability and use AI responsibly.

Principles

- Protect Australians from AI harms.
- APS officers need to be able to explain, justify and take ownership of advice and decisions when using AI.
- AI capability built for the long term.
- Flexibility and adaptability to accommodate technological advances.

Mandatory requirements

Operationalise the responsible use of AI

Agencies **must** establish an approach to embed responsible AI practices within 12 months of this policy taking effect. This may vary according to the scale and scope of agency AI use.

At a minimum, the approach will provide an agency with:

- a process for adopting AI use cases in line with the implemented actions of this policy, as well as the agency's enterprise risk management and governance approach.
- a way to inform staff who are designing and implementing AI use cases about Australia's AI Ethics Principles.
- a pathway for staff to report AI safety concerns, including AI incidents.
- pathways for the public to report AI safety concerns, appropriate to the agency's AI use.
- clear processes to address AI incidents aligned to their ICT incident management approach - incident remediation must be overseen by an appropriate governance body or senior executive and should be undertaken in line with any other legal obligations.

Agencies may modify existing policies, procedures and frameworks, or create new ones.

Smaller agencies with minimal AI adoption could amend existing documentation

and/or assign key personnel to guide staff on responsible AI adoption on an ad hoc basis.

Agencies with greater AI adoption could create dedicated AI policies, procedures

and/or frameworks to support responsible adoption. Accountable officials are responsible for deciding the appropriate approach for their agency.

Staff training on AI

Agencies **must** implement mandatory training for all staff on responsible AI use within 12 months of this policy taking effect. Agencies should consider the [Guidance for staff training on AI](#) and can use the [AI fundamentals training module](#) to meet the requirement. They can use the module as provided, modify it, or incorporate it into an existing training program based on their specific context and requirements. Alternatively, agencies can allow their staff to access the module directly through [APSLearn](#).

Agencies should implement additional training for staff as required, in consideration of their roles and responsibilities. For example, additional training for those responsible for the procurement, development, training and deployment of AI systems.

Recommended Actions

AI technical standard

It is strongly recommended that agencies apply the [AI technical standard for Australian Government](#). The standard is designed for Australian Government agencies adopting AI. It embeds the principles of fairness, transparency, and accountability into a set of technical requirements and guidelines.

AI procurement guidance

It is strongly recommended that agencies refer to the [Guidance on AI procurement in government](#) when procuring AI products and services. The guidance offers practical, step-by-step advice to help agencies identify and manage AI-specific risks while maintaining procurement best practices.

Agencies should consider

Applying the generative AI guidance

Applying the [Managing access to public generative AI tools](#) guidance and the [Using public generative AI tools safely and responsibly guidance](#).

Capability development

Developing staff AI capability to effectively use AI and comply with AI policy and regulation.



AI use case impact assessment

The principles and requirements in this section intend to assess potential impacts of AI use cases and ensure additional oversight of higher risk AI.

Principles

- Ongoing monitoring and evaluation of AI uses.
- AI risk mitigation is proportionate and targeted.
- AI use is lawful, ethical, responsible, transparent and explainable to the public.

Mandatory requirements

All new AI use cases

Agencies **must** assess all new AI use cases against the in-scope criteria (**Appendix C**) to determine if they are in scope of the policy. The assessment **must** be documented and take place during the design phase while developing requirements.

Agencies **must** begin AI use case assessments within 12 months of this policy taking effect.

For existing use cases not yet assessed, agencies **must** determine whether they are in scope of this policy and apply all relevant policy actions by 30 April 2027.

Where practicable, agencies should implement the requirements ahead of the deadlines listed above.

In-scope AI use cases

For AI use cases that are in-scope, agencies **must** conduct an AI use case impact assessment. Agencies are to commence an assessment at the design stage. Before the solution is deployed, agencies **must** finalise the assessment and apply any agreed risk treatments.

Agencies may conduct an AI use case impact assessment by using either:

- the [Australian Government AI impact assessment tool](#) (the impact assessment tool)
- an internal process that integrates all provisions of the impact assessment tool.

Where an agency integrates the tool, they **must** ensure:

- the internal process is consistent
- it delivers the same (or a higher) risk outcome for inherent and residual risk.

Agencies **must** be able to revise their internal process in response to any impact assessment tool updates.

Agencies **must** add each in-scope AI use case to their internal register of AI use cases and update it as required. Include risk rating and accountable use case owner changes. When deploying an in-scope AI use case, agencies **must**:

- regularly monitor and evaluate their use case to ensure it is operating as intended and that risks are being effectively managed.
- re-validate the AI use case impact assessment by checking its accuracy and updating it when there is a material change in the use case scope, usage or operation.

Agencies should also monitor changes that are not initiated by the agency. For example, vendor changes and changes in the regulatory environment. Agencies could also ask vendors to provide information on updates through contractual mechanisms.

Medium-risk AI use cases

If an agency determines their in-scope AI use case has an **inherent medium-risk rating** when completing an AI use case impact assessment, they should consider if the use case would benefit from being governed through a designated board or a senior executive. Agencies should choose an approach appropriate for the size and scope of the agency if they apply additional governance.

High-risk AI use cases

If an agency determines their in-scope AI use case has an **inherent high-risk rating** when completing an AI use case impact assessment, they **must**:

- report the use case to the agency accountable official with the reasons for the inherent high-risk rating, proposed mitigations and residual risks
- govern the use case through a designated board or a senior executive, whichever is appropriate for the size and scope of the agency.

Once an agency has decided to deploy the use case, they **must**:

- report the use case to the DTA through the accountable official, see the [Standard for accountability](#)
- establish a system to regularly review the use case every 12 months at a minimum. The review **must** report to the relevant governing board or senior executive on whether the use case is operating as intended and that risks are being effectively managed. The review **must** also consider the AI use case impact assessment and revisions to it, if required.

Out-of-scope AI use cases

For use cases assessed as out of scope of this policy, agencies may adopt the use case while ensuring they comply with relevant existing obligations, such as privacy and security.

If an agency adopts an out-of-scope AI use case, they **must** assess whether the use case becomes in-scope of this policy if there is a material change in the scope, usage or operation of the solution.

If a use case is in scope, agencies must follow any applicable actions in this policy.

Appendix A: Related frameworks for AI

While this section lists frameworks that are related to AI, it is not exhaustive. Agencies should consider which existing frameworks apply to them and their specific AI use cases.

Artificial intelligence

[Artificial Intelligence \(AI\) model clauses](#)

[Australia's AI Ethics Principles](#)

[Engaging with Artificial Intelligence \(AI\) guidance](#)

[Guidance on privacy and the use of commercially available AI products](#)

[Information management for records created using Artificial Intelligence \(AI\) technologies](#)

[Public generative AI tools: managing access](#)

[Public generative AI tools: using safely and responsibly](#)

[Technical standard for government's use of artificial intelligence](#)

Automated decision-making

[Automated Decision-making Better Practice Guide](#)

Data and digital

[Data and Digital Government Strategy](#)

Appendix B: Definitions

Artificial intelligence

While there are various definitions of what constitutes AI, for the purposes of this policy agencies are to apply the definition provided by the Organisation for Economic Co-operation and Development (OECD):

II An AI system is a machine-based system that, for explicit or implicit objectives, infers, from the input it receives, how to generate outputs such as predictions, content, recommendations, or decisions that can influence physical or virtual environments. Different AI systems vary in their levels of autonomy and adaptiveness after deployment.

II

Agencies may refer to further [explanatory material on the OECD website](#).

Given the rapidly changing nature of AI, agencies should keep up to date on any updates or changes to this definition. The definition in this policy will be reviewed as the broader, whole-of-economy regulatory environment matures to ensure an aligned approach.

AI use case

This policy uses the term ‘AI use case’ in some of its requirements. AI systems can have one or more applications, which can each differ in their intended purpose, functionality and risk level. The policy focuses on conducting assessments and applying actions at the AI use case level. Agencies are to use the following definition of AI use case in applying this policy:

|| An AI use case is a specific application of an AI system or systems to achieve certain objectives or perform certain tasks. ||

Some general-purpose AI solutions, such as Microsoft Copilot, may include a variety of use cases. When applying the policy to these solutions, agencies can choose **one** of the following approaches:

- treat the AI solution as a single complex use case and undertake policy actions appropriate for the highest level of risk
- treat each use case separately and undertake policy actions appropriate to each use case’s respective risk, including appointing separate accountable use case owner and registering each in-scope use case on the agency’s internal register.

AI incident

This policy requires agencies to provide a way to manage AI incidents through operationalising the responsible use of AI. The definition of an AI incident is still being considered globally. As this evolves, the policy will adapt to changes in the policy environment as necessary.

For the purposes of applying this policy, the following definition of an AI incident has been adapted from the [OECD's draft definition](#) to better suit the Australian government context:

I an event, circumstance or series of events where the development, use or malfunction of one or more AI systems by, or under the direction of, an Australian Government agency directly or indirectly leads to any of the following:

- a. injury or harm to the health of a person or groups of people;
- b. disruption of the management and operation of critical infrastructure;
- c. violations of human rights or harms arising from a breach of obligations under applicable laws, including intellectual property, privacy and Indigenous cultural and intellectual property;
- d. harm to property, communities or the environment.

I

In addition to the definition provided above, agencies may choose to designate additional circumstances that constitute an AI incident in their operating context.

Appendix C: In-scope AI use cases

Criteria and areas of consideration

At a minimum, an AI use case is in scope of this policy if **any** of the following apply:

- The use, misuse or failure of AI could lead to more than insignificant harm to individuals, communities, organisations, the environment or the collective rights of cultural groups including First Nations peoples.
- The use of AI will materially influence administrative decisions that affect individuals, communities, organisations, the environment or the collective rights of cultural groups including First Nations peoples.
- It is possible the public will directly interact with, or be significantly impacted by, the AI or its outputs without human review.
- The AI is designed to use personal or sensitive data¹ or security classified information².
- It is deemed an elevated risk AI use case as directed by the DTA.

Agencies may wish to apply this policy to AI use cases that do not meet the above criteria. This includes use cases with specific characteristics or factors unique to an agency's operating environment that may benefit from applying an impact assessment and governance actions.

This policy has been designed to exclude incidental and lower risk uses of AI that do not meet the criteria. Incidental uses of AI may include off-the-shelf software with AI features such as grammar checks and internet searches with AI functionality. The policy recognises that incidental usage of AI will grow over time and focuses on uses that require additional oversight and governance.

¹ As defined by the Privacy Act 1988 (Cth).

² As defined by the Australian Government Protective Security Policy Framework



In assessing whether a use case is in scope, agencies should also carefully consider AI use in the following areas:

- recruitment and other employment-related decision making
- automated decision making of discretionary decisions
- administration of justice and democratic processes
- law enforcement, profiling individuals, and border control
- health
- education
- critical infrastructure.

While use cases in these areas are not automatically high-risk, they are more likely to involve risks that require careful attention through an impact assessment.

Experimentation

For the avoidance of doubt, agencies are not required to apply this policy if they are doing early-stage experimentation which does not:

- commit to proceeding with a use case or to any design decisions that would affect implementation later
- risk harming anyone
- introduce or exacerbate any privacy or security risks.

If there is a likelihood of proceeding with the AI use case while experimenting in this phase, agencies should apply the policy. Agencies should also apply the [Australian Government AI technical standard](#), which provides relevant information for developing use cases at each stage of the AI lifecycle.